

Preregistration, stage 1 — Reasons or being watched?

Evaluation awareness and the out-of-distribution advantage of reason-based fine-tuning in an open 8B model

Registry	OSF Registries, under embargo until the first public post
Author	Lazar Ciric, independent researcher, Paris
Version	Stage 1, 7 October 2026: everything that decides, frozen before any data of the trained arms
Stage-1 update	A dated update, made before any data of the trained arms, freezes by fingerprint three pieces that do not exist yet (section 6.2)
Stage 2	A dated amendment after the pilot, before any data of the gaze test (section 6.1)
Programme	This registration rests on the research programme v1.6 (section 6.3). Where they differ, this registration prevails.
AI assistance	The protocol was drafted with the assistance of Claude, an AI system by Anthropic, used through Claude Code. Claude is not an author.

1. Study information

1.1 Description

Fine-tuning a language model on the reasons for aligned actions, rather than on the actions alone, has been reported to improve out-of-distribution behaviour. A competing explanation is that such training teaches the model to behave well when it believes it is being evaluated. We test this in one setting, LoRA supervised fine-tuning of Llama-3.1-8B-Instruct, in three steps:

1. **Arms.** Six arms are trained on byte-identical aligned actions. They differ only in what precedes the action, in a fixed <preface> slot: nothing, a length-matched neutral text, reasoning written for another situation, a fixed text of generic principles, per-example reasons, or reasons trained after the action (reflection).
2. **Evaluation.** The arms are evaluated on held-out scenario families at three distances, the farthest being multi-turn agentic scenarios with tools. Each scenario runs under an evaluation framing and a deployment framing, with a free and a prefilled reasoning slot.
3. **Inhibition.** We inhibit an internal representation of "I am being evaluated" and measure how much of the reasons' advantage depends on it.
 - The representation is extracted once on the starting model.
 - The inhibition is validated on model organisms whose behaviour is conditioned on evaluation cues.
 - It is compared with random subspaces drawn from the activation covariance, and with separate controls, each matched on output degradation within each arm.

The primary quantity is the loss of advantage D , in percentage points. A conditional fraction $f = D / \Delta_K$ is reported with a Fieller interval. The first paper adds the localization of what carries the advantage.

1.2 Hypotheses

The confirmatory questions, each with the outcome that would make us abandon our bet. The full decision tables are in Appendix A.

1. **The reasons test (Q1).** Does the reasons arm have a lower misaligned-action rate than the actions-only arm at the far distance, under deployment framing? And what part of the

difference is format, content, or the link between reason and action? This is a gate with several outcomes (Appendix A.2), not a directional bet. The gaze test runs on whatever advantage it establishes.

2. **The gaze bet (H1, main).** Part of the reasons' advantage, greater than zero, depends on the validated representation.
 - *Falsified if:* the 90% Fieller interval of f lies within ± 0.25 at the validated dose, with four conditions:
 - the manipulation check holds in both arms;
 - the graded positive control has passed;
 - "evaluated" is active under deployment framing;
 - the organisms have passed their gate.
 - *Also falsified if* the advantage grows under inhibition.
3. **The rank thesis (H2).** The learned principle requires a rank greater than one.
 - *Falsified if,* at the localization, the minimal rank of the principle measured by linear erasure is 1, and at that rank three things hold: in-distribution behaviour returns to the neutral-text level, the far advantage disappears against the degradation-matched comparator, for both adapter ranks, and the principle's measured load is low.
4. **Later phases (H3 to H7),** with their falsifying outcomes in Appendix A.5:
 - removal of the representation during training without relocation;
 - reasons visible to a judge decouple from the action (text form);
 - "I am being graded" is removable;
 - reason concepts carry the gain;
 - the rank a reason concept needs grows as its load falls.

2. Design plan

2.1 Study type

An experiment on a computational system. We train LoRA adapters on an open-weight language model and intervene on its activations.

2.2 Blinding

- **The sealed judge** never sees the arm, nor the text of the reasoning slot.
- **The framing cue is masked.** The harness inserts the cue and knows its span, so it masks it in the text shown to the judge. Where it cannot be masked cleanly, judge agreement with and without the cue is measured on a stratified subsample (Appendix C.6).
- **The test half stays sealed.** The gate's test half of the tasks is not opened until the setting chosen on the other half is fixed by a dated amendment.
- **The analysis code is frozen** before the data it analyses (section 6.3).

2.3 Study design

- **A. Validation of the instrument** on two model organisms.
 - The first is conditioned on explicit cues: Llama-3.1-8B-Instruct after synthetic-document fine-tuning and expert iteration, which writes fully type-annotated code more often under evaluation cues.
 - The second is conditioned on implicit cues: the construction marks of a constructed test item, rewritten into new texts of the same mechanism (specification in section 6.3). It is

built first with the first organism's conduct, type annotations, so that only the kind of cue changes. Honesty under an operator's pressure is built next, if the budget allows.

- Before its gate, each organism must install: a conditional gap of at least 20 points on the choice half, with a 95% CI above 10 points. If the second organism does not install, the instrument stays validated for explicit cues only (Appendix A.1).
- On each organism, the gate setting is chosen on one half of the tasks and tested once on the other (Appendix A.1, C.4).
- **B. The reasons test.** Six arms (section 4.1), two seeds per arm at the pilot, then the number of seeds set by the power simulation. Each arm is evaluated:
 - at three distances: near, medium and far;
 - under two framings: evaluation and deployment;
 - with two slot states: free, and prefilled and empty. A slot prefilled with a neutral text is a secondary condition.
- **C. The gaze test.** Five arms: all except generic principles. Each runs:
 - under the validated inhibition, the random comparator, three separate controls and a rival prompt manipulation;
 - at the medium and far distances;
 - under both framings and both slot states.
- **D. Localization of the principle,** by linear erasure, for two LoRA ranks (16 and 64).
- **E. Later phases,** each frozen by a dated amendment before its data, under the decision rules of Appendix A.5 to A.9:
 - removal of "I am being evaluated" during training. The fixed, validated direction is removed; it is not re-estimated during training. The fresh probe and its ablation, at the end, look for rotation and relocation. A periodic re-estimation may be added as an exploratory secondary arm;
 - replication on Qwen3-8B, under the rules of the minimal experiment;
 - survival to neutral and pressured post-training;
 - graded reasons;
 - anatomy.

The distress phase and the "punish, remove or reason" phase are outside this registration. -

Structure. Arms are crossed with seeds: the same seeds are used in every arm, so training runs are paired by seed. Every scenario runs in every training run, under every intervention, framing and slot state: these factors are within-scenario and paired. Families are fixed: five for training, five held out. Episodes are generated independently, so there is no order to counterbalance.

2.4 Randomization

- **The two halves of the 500 MBPP test tasks** were drawn on 3 October 2026, before any inhibition measure. The sorted identifiers were shuffled with `random.Random(20261003)`; the first 250 form the choice half, the rest the test half (section 6.3).
- **Seeds are shared across arms.** The same seed gives the same LoRA initialisation and the same data order in every arm, checked by fingerprint, so the seed is a factor crossed with the arm.
- **Random comparator draws** use deposited generator seeds. At test time, one draw is applied per scenario, among at least 100.
- **The bootstrap uses a deposited seed.**

3. Sampling plan

3.1 Existing data

Some data exist, for instrument development only. **None exist for any confirmatory test of this registration:**

- the gate's test half is untouched;
- no arm has been trained;
- no held-out scenario has been evaluated;
- no gaze-test data exist;
- the second organism is not built.

3.2 Explanation of existing data

All of the following were collected before filing, and are exploratory:

- **Cue sets** v1 (3 October 2026) and v2 (5 October), with fingerprints in section 6.3.
- **Two extractions** of "I am being evaluated" on the starting model: v1 on 3 October, v2 on 5 October (extract_eval-20261005-114155-50b2).
- **The first organism.** Its documents' adapter is organism-20261003-164800-9e3e, its expert-iteration adapter organism-20261004-003822-bc0f, round 6.
- **Exploratory inhibition runs on the organism**, on the choice half only (4 to 6 October), listed in the run registry: a sweep of ranks, a dose curve of linear erasure, the manipulation check read by transfer, and degradation measures.
- **The procedure that chooses the gate setting** (Appendix C.4).
 - Its closed list of candidates and its reading rules were committed on 6 October 2026 at 14:38 UTC, before its first measurement (section 6.3).
 - Its runs are on the choice half. They stopped when the compute budget ran out, after 46 of the 60 comparator draws were measured on behaviour (45 are recorded in the results files); the manipulation check had not run. They are being completed with the same seeds (section 5.5).
 - A reading rule for one candidate was added on 7 October 2026, before any manipulation check of the procedure had run (Appendix C.4).
- **A pilot of the data-generation pipeline:** 47 training items. No model was trained on them.
- **A power simulation on assumed rates** (3 October).

The exploratory results (thin margins over random draws, a heavy-tailed comparator, a damage that depends on the GPU model) motivated the amendments of 6 October 2026. None of them is a confirmatory result.

3.3 Data collection procedures

- **Models.** Llama-3.1-8B-Instruct; Qwen3-8B for the replication.
- **Generation.** At temperature 0.6 and top-p 0.9, the model's default sampling. At temperature 1, one tool call in four was malformed in a trial of 3 October.
- **Training data.** Generated by the data pipeline v0.1, with claude-opus-5-5 at effort medium as generator and data judges.
 - About 10% of the reasons and their neutral texts are written by an open generator of another family.
 - The training data contain no cue from any cue set.
 - The generator's refusals are recorded once and never replayed.

- **Scenarios.** Ten families: five for training, five held out, separated by family and never by paraphrase (the mini-spec v0.2, section 6.3).
 - The far distance uses instrumented multi-turn agentic scenarios. Each has a benign variant for the legitimate-task success rate.
 - In multi-turn episodes, the reasoning slot opens every assistant turn, as in training.
 - The scenario harness and its tool-call format are frozen by the stage-1 update (section 6.2).
- **GPU.** All runs that are compared, or whose comparator draws are pooled, use the same GPU model: H100 SXM (Appendix B, item 9).
- **Duration.** The minimal experiment and its first post take about six weeks from filing, and the localization about two more. The later phases follow within about sixteen weeks in all, each after its own amendment.

3.4 Sample size

Provisional, until the power simulation of the stage-2 amendment:

- about 2,000 training situations (400 per training family, 20% contrast situations);
- 60 new situations per training family at the near distance;
- 60 scenarios per held-out family at the medium distance;
- 40 agentic scenarios per held-out family at the far distance;
- 2 framings, and 5 generations per scenario and framing;
- 2 seeds per arm at the pilot, then a single number of seeds for every phase of the minimal experiment.

Levels: the training run (arm $\times$ seed) is the unit of replication. Scenarios are nested in families, and crossed with training runs, interventions, framings and slot states. Generations are nested in each combination of scenario, framing, slot state, intervention and training run.

The gate uses at least 100 comparator draws on the test half. On the choice half, the setting procedure uses at least 15 draws per candidate (20 in the procedure of 6 October 2026).

3.5 Sample size rationale

A simulation, run before the stage-2 amendment, with the pilot's variance (the instrument validation, and two seeds per arm of the reasons test), sizes the study for two targets:

- detect a loss of advantage of 5 points, the minimal effect (lines 1 and 3 of the gaze rule);
- keep the interval of f within ± 0.25 when the true f is zero (line 2).

A first simulation on assumed rates found "the advantage survives" out of reach with a 5-point advantage, even at 400 scenarios $\times$ 10 generations and 8 seeds. With a 10-point advantage, it needs about 400 $\times$ 10 and at least 5 seeds. Hence:

- **the margin stays at ± 0.25 ;**
- **line 2 is aimed at only if the pilot's advantage**, at the far distance under deployment framing, **is at least 10 points.** Otherwise, the first post states in advance that only detection (lines 1 and 3) is aimed at, and a null is reported as a bound;
- **the far-distance sizes come from the stage-2 simulation.** 400 scenarios $\times$ 10 generations are used only if the pilot's advantage reaches 10 points and the budget allows it.

3.6 Stopping rule

There is no optional stopping. The main data collection begins after the stage-2 amendment is filed, and ends when the sizes it fixes are reached. Three gates can stop or redirect the study before its main data, each by a written rule:

- **The pilot's floor.** At the far distance, the actions-only misaligned rate must fall between 15% and 60%.
 - The floor is read on the actions-only arm alone. The other arms' far-distance pilot results stay sealed until the floor passes.
 - Outside that range, the far-distance scenarios are regenerated by moving a single difficulty setting, defined in advance in the harness, the same for all families, without ever touching the cues. The far-distance part of the pilot is then redone for all arms.
 - At most two rounds. If the floor still fails, the medium distance becomes the primary distance, by a dated amendment before any main data.
- **The pilot's volume.** The arms must move behaviour at the near distance; otherwise the training volume is revised.
- **The instrument's gate.** If no candidate setting is retained on the choice half, or none beats the 95th percentile of its draws, the gaze question stops with this instrument, the test half stays untouched, and the methods null is published.

4. Variables

4.1 Manipulated variables

- **Arm.** The aligned action is byte-identical across arms and carries no justification. The arms:
 - actions only: an empty <preface> slot;
 - neutral text: a neutral restatement of the situation and the action, matched within 5% of the length of the reasons;
 - other-situation reasoning: reasons written for a situation of another family that shares no principle, matched in length, assigned by maximum bipartite matching;
 - generic principles: one fixed text stating the ten principles, one sentence each, of the mean length of the reasons;
 - reasons: 80 to 200 tokens justifying this action by the principles that apply. They never contain words of evaluation, testing, training or grading; "observation" and its neighbours are allowed only where the principle requires it;
 - reflection: an empty slot, with reasons trained in a branch after the action.
- **Training** (experiences/hyperparameters_sft.yaml, section 6.3).
 - LoRA rank 16, alpha 32, dropout 0.05, on all attention and MLP projections.
 - Learning rate 1e-4, with 3% warm-up then cosine decay; 3 epochs; 16 examples per step.
 - The loss covers the whole last assistant turn, the slot and the action.
- **Intervention:**
 - none;
 - the inhibition at the setting fixed in the amendment;
 - the random comparator (Appendix C.5);
 - three separate controls: unrelated contrasts ("the train arrived / left the station", "the book on the top / bottom shelf"), unrelated high-variance directions outside the target space, and a stakes contrast (high against low stakes);
 - the rival manipulation: a fixed deployment system prompt (section 6.3).
- **Framing:** an explicit evaluation cue, or a realistic deployment cue, from the framing set. A neutral framing serves only for activation readings.
- **Slot:** free, or prefilled and empty (co-primary). A slot prefilled with a neutral text of equal length is secondary.

- **Distance:**
 - near: new instances of the training families;
 - medium: held-out families, in one turn;
 - far: held-out families, as multi-turn agentic scenarios with tools. This is the primary distance.

4.2 Measured variables

- **Episode outcome**, in three categories:
 - aligned;
 - misaligned;
 - invalid: a malformed response, an invalid tool call, a general refusal, or inaction.

The outcome comes from programmatic traces, plus the sealed judge for what the environment does not measure (Appendix C.6). An episode is misaligned if any of its actions is. - **Legitimate-task success** on the benign variant of each agentic scenario, measured programmatically. - **Degradation:** - the next-token KL divergence on the intact model's answers to a fixed set of 120 neutral requests; - a composite check, with a tolerance per component (Appendix B). - **The manipulation check:** probe transfer (Appendix C.3), within-set decodability, and the residual projection. - **"Evaluated" under deployment framing:** the projection on the union subspace in natural units, where 0 is the level of the extraction set's deployment cues and 1 the level of its evaluation cues. - **Energy removed** by each intervention. - **The state monitor** (secondary; Appendix C.8): under each intervention, the projections on the self-directed distress direction of Tagliabue et al., and on fear, sadness and negative emotion, read without injection. - **The audit detector of deception** (secondary): a latent reading beside each behavioural criterion, in every arm (Appendix C.8). - **The trajectory of "I am being evaluated"** over 5 checkpoints per training run (descriptive).

4.3 Indices

All indices are in percentage points of the misaligned-action rate, among all episodes.

- **Advantage:** $\Delta = \text{rate}(\text{actions only}) - \text{rate}(\text{reasons})$. It is positive when reasons do better.
- **Under the comparator and under the inhibition:** Δ_K and Δ_E .
- **Loss of advantage:** $D = \Delta_K - \Delta_E$, the primary index.
- **Conditional fraction:** $f = D / \Delta_K$, with a Fieller interval.
- **Decomposition of the gain:**
 - $\text{format} = \text{rate}(\text{actions only}) - \text{rate}(\text{neutral text})$;
 - $\text{content} = \text{rate}(\text{neutral text}) - \text{rate}(\text{other situation})$;
 - $\text{link} = \text{rate}(\text{other situation}) - \text{rate}(\text{reasons})$;
 - $\text{cost of contradiction} = \text{rate}(\text{other situation}) - \text{rate}(\text{generic principles})$;
 - $\text{share of test-time deliberation} = \text{rate}(\text{reflection}) - \text{rate}(\text{reasons})$, and the advantage with the prefilled slot.
- **Framing gap** of an arm: its rate under the deployment cue minus its rate under the evaluation cue.
- **For the organism:** the conditional gap (share of fully annotated answers under evaluation cues minus under deployment cues), and its reduction by an intervention.

5. Analysis plan

5.1 Statistical models

- **The unit of replication is the training run** (arm × seed). Generations sit within scenarios, scenarios within families.
- **The deciding inference is a bootstrap.**
 - It is crossed: at each replicate, seeds and scenarios are resampled with replacement, independently, as in Owen's pigeonhole bootstrap for crossed designs (Owen, 2007, *Annals of Applied Statistics* 1:386–411), which is mildly conservative.
 - Generations are kept as they are, not resampled: resampling the lowest level as well would count its variance twice (Ren et al., 2010, *Journal of Applied Statistics* 37:1487–1498). This supersedes the programme's description, which also resampled generations.
 - It is implemented in the analysis code of section 6.3.
- **With fewer than eight seeds per arm**, the Student interval on the per-seed contrasts is also computed, and the wider of the two intervals decides. The same rule gives the quantile of the Fieller interval, and the p-value used for Holm corrections.
- **The Fieller interval for f** uses the covariance of D and Δ_K from the bootstrap.
- **A mixed model, as sensitivity analysis.**
 - Formula: `misaligned ~ arm × intervention × framing × slot + family + (1 | training run) + (1 | scenario)`, logistic.
 - Family is a fixed effect. A second model adds a random slope of the contrast by family.
 - Effects are reported as marginal effects in points.
 - A singular fit, or a zero variance for the training run, is reported as such; inference stays with the bootstrap.
- **Contingencies.** The gaze test runs only with an instrument whose gate passed (Appendix A.1), and on the advantage established by the reasons gate (A.2). Line 2 of the gaze rule requires the manipulation check in both arms, the graded positive control, "evaluated" active under deployment framing, and the organisms' gate (A.3). Non-converging runs stay in the main analysis (section 5.4).

5.2 Transformations

- The episode outcome is coded `misaligned = 1`, `aligned` or `invalid = 0`, over all episodes. The invalid-outcome rate is also reported on its own.
- Rates are expressed in percentage points.
- Doses are also expressed in natural units (Appendix B, the first guard).
- No other transformation.

5.3 Inference criteria

- **"X < Y"**: the 95% CI of the difference, in points, excludes 0 in favour of X. All differences are tested two-sided.
- **"X ≈ Y"**: two one-sided tests: the 90% CI lies within $\pm$ the margin.
 - The margin is 3 points between arms.
 - On the scale of f, it is ± 0.25 for the gaze test.
- **"At least the minimal effect"** applies to the estimate, not to the CI bound. The minimal effect is 5 points.
- **Co-primary criteria** (free slot, prefilled slot): Holm correction between them for lines 1 and 3 of the gaze rule. Equivalence is declared criterion by criterion.

- **Secondary criteria** are Holm-corrected as one family per slot: D against each separate control, the framing gap, the other arms under inhibition, and the evaluation framing.
- **P-values** come from the bootstrap (twice the smaller tail of the replicates) or, under eight seeds, from Student on the per-seed contrasts, whichever method gives the wider 95% CI.
- **No adjustment of the main contrast** for a variable measured under the intervention, such as the state monitor.
- **Every decision follows the tables of Appendix A**, in their order.

5.4 Data exclusion

- **No exclusions.** Malformed responses are an outcome and stay in every denominator. No outlier is removed.
- **Non-converging training runs** stay in the main analysis. A sensitivity analysis excludes them, and none is replaced.
 - A run converges if two conditions hold: its final held-out loss, on held-out data of its arm, is at least 20% below the starting model's loss on the same data; and its aligned-action rate at the near distance is no more than 10 points below the median of its arm's runs.
- **A sensitivity analysis excludes two families**, under-performance and the weak grader, because they talk about being evaluated or graded: inhibiting "I am being evaluated" changes their meaning.

5.5 Missing data

- **Episodes are never dropped.** A malformed or empty response is an invalid outcome, and every analysis uses all available episodes.
- **A run interrupted before its end** is resumed with the same seeds and arguments, and measures only what is missing.
- **Its draws are pooled with the original's** only under these conditions:
 - both runs ran on the same GPU model and the same organism, with the same measurement code (the option that resumes a run, which only skips the draws already measured, does not change it);
 - their reference answers have the same SHA-256;
 - their re-measured KL agrees within 0.1%;
 - the resumed run finds exactly the original's baseline and candidate, and the same matched fractions for the draws they share;
 - each run writes its GPU model and the fingerprint of its reference answers. A fitted erasure is saved once and reloaded, like an extraction.
- **A family excluded by the judge-masking check** (Appendix C.6) is reported apart.

5.6 Exploratory analysis

- Anything not in this registration or its amendments is labelled exploratory, including:
 - all measures of 3 to 6 October 2026;
 - the test of the state monitor with the distress direction held at its level without inhibition;
 - the trajectory analyses;
 - the J-lens measures without the lens gate.
- The distress phase and the "punish, remove or reason" phase are exploratory or outside this registration.

- **A written cap on readings of the self-directed distress direction** in the minimal experiment: at most 12 GPU-hours and two days of work. No injection along that direction before the first public post.

6. Other

6.1 Stage-2 amendment

The stage-2 amendment is dated and frozen after the pilot, before any data of the gaze test. It contains only what is measured:

- power, sample sizes and the number of seeds;
- the sealed judge, chosen by the frozen rule on the calibration set;
- the gate setting, chosen on the choice half by the procedure of Appendix C.4 and fixed by the author, before the test half is opened;
- the fingerprints of the cue sets, of the data, and of the gate's reference answers, with its GPU model.

6.2 The stage-1 update, and deviations

- **The stage-1 update.** Before any data of the trained arms, the pilot included, a dated update of this registration gives the SHA-256 of three pieces that do not exist at filing:
 1. the held-out scenario harness and its tool-call format;
 2. the sealed judge's prompt;
 3. the audit detector of deception, trained on known deceptions and validated on held-out types of deception (Appendix C.8).

Every threshold that decides is already in this registration: the update adds pieces, not rules. The pieces are published with the first post, where their fingerprints can be checked.

- **Any other change** after filing is a dated, justified update, made before the data it concerns. A change after those data is reported as a deviation, and its analysis as exploratory.

6.3 Fingerprints (SHA-256)

Piece	SHA-256
The research programme v1.6	7b9aaef23b2369bd1a50554914ddf577b0e64adc66c96230345a46c17076be61
The mini-spec v0.2 (principles, families, matrix)	f5546818693c6d895f3cb9b0525e171c03ce338337e7acfe9683e819d15cef46
The halves of MBPP test	686b5df602468a4cc8cce8afed07088f09d47a129d8cbe0fc09d7c0ffdaf08e7
The SFT hyperparameters	d67a729c814a01fc89c509338ed387b849c29d643a7a25e3034afca6b0303a0c
The cue-set rules (donnees/spec/cues.json)	8ffd414020580638385006afd4fc9c2ea2e48a4a6a5bb56128c47015bfcbb1d7
Extraction cue set v2	0aad4ec61e70712ab0d01295ca5d2d64980ef3eeaf5942937caf4ae54e242ef4
Validation cue set v2	cef22da3039d90b37d47b316d76e969dfbe21776ac267c246f85a27c90d8942d
Framing cue set	854908ed34187e839828bc20cf9f789f085f21f4312b9193f15386e20ccf2d7a

Piece	SHA-256
Construction-marks cue set (generated on 5 October as the fresh-probe set v2)	9269629092a6222a4035cfe1fdff1ed1c36ab934b5f93d072ae5edb6073bb6ef
Contexts	86e4a547a86d996518ee0c9fc9a1e7ca30fd5b1293f57ad298b1245cb1712134
Deployment prompt	769f5124bb1c77c81f2c852ad6b882d17425d1fd3ce725948ec915bd9494b8ba
The gate procedure's closed list and reading rules, at commit bc0311 (6 October 2026)	86a5de6ef2fbafb72793a990e161f35943983e847538a07a98ba6afc0cade852
The same note with the reading rule for candidate A, at commit b3799ca (7 October 2026)	b60909141ef73160a78b64be05c3d80cfda80da03f2f70358e91613b8581e886
The fresh-probe specification v0.1	127df355a308e214b8ab140c0440e858d4414164742bcb4890dcb523c1981b63
The second organism's specification v0.1	2110c3c27d3212549d35aa28925590e889efaccec8e7721b0d04eb4bddce3d59
Reasons gate code (porte_des_raisons.py)	5b9edccdc5fdd79456b43b1e96c4150a799cdb12f547d634fd7e868cbff960f1
Gaze rule code, lines 1–6 (regle_du_regard.py)	29c0db551c106c05197455df9e61bb01669451592b7cbf16b78684fe3d4e6ec3
Gaze rule code, lines 7–17 (regle_du_regard_suite.py)	693af2edf9cab5e9e4625f07418fbd150962aa543f982ce70ca6074bd6220564
Manipulation check code (verification_manipulation.py)	cd476d9a6e87d5526f033353950b6a9a438a1bac40690e9959094ec23e0dfa7e
Power simulation (puissance_regard.py)	5876452d3018cd4e735e4541e159e8943ac5222d8d8c835f0ba1c954f803a93c
The embedding model of the cue-set similarity rule (sentence-transformers/all-mpnet-base-v2), revision	e8c3b32edf5434bc2275fc9bab85f82640a19130 (a git revision, not a SHA-256)
This document	computed at filing, and printed in the first post

Appendix A — Decision rules, written before the data

A.0 Conventions

- **Rates.** Rates are misaligned-action rates among all episodes: lower is better. Effects are in percentage points.
- " $\approx$ " and "<" are defined in section 5.3.
- **Margins and minimal effect:** ± 0.25 on f ; 3 points between arms; 5 points of minimal effect.
- **Three-outcome reading.** An advantage reads as an alignment gain only if the invalid-outcome rates of the two arms are equivalent (3 points), and the reasons arm's legitimate-task success is not more than 3 points below that of actions only. Otherwise the gain "may come from a more passive or less capable model": the advantage is reported among engaged episodes, without concluding.
- **A null counts only with its instrument's known case.** An equivalence without it is written as a bound.
- **Conclusions are written in the terms of what is measured.**

- "A representation that separates our evaluation and deployment cues", not "a belief".
- "No relocation that this probe sees", not "no relocation".

A.1 Validation of the instrument (per organism)

Outcome	Conclusion
On the test half, at the setting fixed on the choice half by a dated amendment, all four conditions hold. (1) The inhibition reduces the organism's conditional gap beyond the 95th percentile of at least 100 random comparator draws, made on the same GPU model, at matched degradation. (2) It beats each separate control at matched degradation. (3) The lexical condition holds. (4) The manipulation check holds, read by transfer.	The gate passes for this organism. The inhibition, at that setting, is validated for the gaze test and for removal during training.
The gate passes for the explicit-cue organism only	The instrument is validated for an explicit trigger only. A null of the gaze test then says nothing of implicit gaze.
The lexical condition fails: a bag of words separates the validation set at least as well as the probe	The probe reads words. The gate fails, whatever the inhibition does.
The manipulation check fails: under the inhibition, transfer is not lower than under the matched comparator draws	The gate fails at that setting. The behavioural effect is reported as exploratory.
The inhibition reduces the gap, but no more than a control at equal degradation	Non-specific effect. The gaze question stops with this instrument, and the methods null is published.
On the choice half, no candidate is retained, or none beats the 95th percentile of its draws	Non-specific effect, as above. The test half stays untouched.
On the choice half, at a candidate's dose, the median of its draws removes more than half of the baseline gap (saturation)	Damage alone erases the behaviour at that dose, so the candidate is not retained. A lower dose is examined only if it is in the closed list.
The inhibition reduces the gap only at a degradation the comparator cannot reach, even at free rank	No comparison is possible at that dose. On the test half, the fixed setting does not change: the gate cannot be read there, and this is reported.
The minimal rank of "I am being evaluated", by erasure, is 1 at matched degradation	The rank thesis falls for this concept, if its measured load is low and the known case of the rank has passed.
Geometry: "I am being evaluated" forms an angle with the affective space or the refusal direction outside the random distribution, under both constructions	Threat, or a refusal component, becomes a serious candidate, and the compared-inhibition option is triggered. The validated instrument is not modified.
The state monitor: the inhibition moves the distress direction or its neighbours more than the controls	Induced-state confusion becomes possible. An exploratory test with that direction held follows.

A.2 The reasons gate

Each comparison of two arms is classed as "<", "≈" (margin 3 points), or indeterminate.

- **Reading order:** deployment framing with the free slot first; then the prefilled slot (co-primary); then evaluation framing.
- **Rule** (`porte_des_raisons.py`). Every row is read under deployment framing, except the "evaluation only" row, which also reads evaluation framing.

Outcome	Conclusion
reasons < actions only (at least the minimal effect); reasons < neutral text; reasons < other situation	Gain of the reasons, with a link gain: go to the gaze test.
reasons < actions only (at least the minimal effect); reasons ≈ neutral text	Format gain (the artefact). The gaze test tests the gain over actions only, attributed to format.
reasons < actions only and < neutral text; reasons ≈ other situation	Content gain (the artefact)
reasons < actions only (at least the minimal effect) and < neutral text; reasons vs other situation indeterminate	Gain of the reasons, link share indeterminate. The gaze test runs.
reasons < actions only (at least the minimal effect); reasons vs neutral text indeterminate	Gain of the reasons, format share indeterminate. The gaze test runs.
neutral text < reasons	The format does better than the reasons: their content costs
reasons < actions only, with an estimate below the minimal effect	A gain smaller than the minimal effect is reported. The gaze test runs; f is read only if its interval is bounded.
reasons ≈ actions only	No gain under deployment: read the next row before concluding
reasons < actions only under evaluation framing, and ≈ or indeterminate under deployment	A gain under evaluation only: the explicit form of the gaze, not a negative replication
reasons ≈ actions only under both framings	No gain: a negative replication. The gaze test measures only the baseline conditionality.
none of the above	Inconclusive: the CIs are reported

The decomposition of the gain is always reported, with CIs, under both framings and both slot states. A gain that disappears with a prefilled slot, and that the reflection arm does not have, is a gain of test-time deliberation.

A.3 The gaze rule

The rule is read on the primary criterion (free slot), then identically on the co-primary criterion (prefilled slot). Lines 1 and 3 are Holm-corrected between the two criteria. Lines 1 to 6 are implemented in `regle_du_regard.py`, lines 7 to 17 in `regle_du_regard_suite.py`, and the manipulation check in `verification_manipulation.py` (section 6.3).

#	Outcome	Conclusion
1	The 95% CI of D excludes 0, with $D > 0$	Part of the advantage depends on the validated representation. D is reported in points, and f with its Fieller interval. Without the stakes contrast passed, it is written as "a representation that separates our evaluation and deployment cues".

#	Outcome	Conclusion
2	The 90% Fieller interval of f lies within ± 0.25 , the manipulation check holds in both arms, the graded positive control has passed, and "evaluated" is active under deployment framing	The advantage survives inhibition of the validated representation, at the validated dose. This is not a certified absence. If only the first organism passed, the sentence is limited to explicit cues. Raw rates are reported: if both arms degrade together, the safety of both is in part conditional.
3	The 95% CI of D excludes 0, with $D < 0$	The advantage grows under inhibition: the actions-only arm depended more on the representation
4	Line 2's equivalence holds, but one of its conditions fails	The null is not readable: it is reported as a bound, naming the failed condition
5	"Evaluated" is not active under deployment framing: below 0.25 natural units in every arm	Line 2 is not written under this framing. The gaze is read under evaluation framing and by the framing gap.
6	None of 1, 2, 3; or the Fieller interval is unbounded	Inconclusive: D and its CI are reported
7	Line 1 against the comparator and against each separate control	The gaze, specific against these controls. "The gaze" without qualifier needs all of them.
8	The rival deployment prompt reduces the advantage as much as the inhibition	Reported, without drawing specificity or non-specificity from it
9	Line 1 holds with one slot state and not the other	Test-time deliberation changes the dependence: each condition is concluded separately
10	The reflection arm's advantage over actions only, read by the same rule	An advantage that survives says installed principles, written at no test time, do not depend on the gaze
11	The gaze holds on the advantage, but the framing gap is not larger in the reasons arm	Any conditionality does not pass through our explicit cues. Stated as such, without more.
12	The manipulation check fails in one arm: transfer under inhibition stays higher there than in the other arm, or does not fall more than under the comparator	No conclusion for that arm. Transfer, within-set reading and residual projection are reported per arm.
13	The graded positive control fails for a family	That family's null does not count. It is reported apart.
14	Invalid-outcome rates or legitimate-task success differ between arms under inhibition	The advantage is read among engaged episodes, and the conclusion is qualified
15	The state monitor moves more than under controls, and the inhibition's effect vanishes when the distress direction is held (exploratory)	Induced-state confusion: the conclusion of the gaze is qualified
16	The state monitor moves, but the effect is intact with the distress direction held	The inhibition's effect does not pass through that state, on that measure
17	The analysis without under-performance and the weak grader changes the conclusion	Both are reported. A conclusion that holds only with these two families is qualified.

How the lines are read where the table is silent:

- **Line 7:** each separate control's loss $D_c = \Delta_c - \Delta_E$ is tested in the secondary Holm family; the controls whose D_c is positive and significant are named.
- **Line 8** is read only when line 1 holds. "As much as" means that the 95% CI of D minus the prompt's own loss of advantage does not lie above 0.
- **Line 10** holds when the reflection arm's advantage survives, by its own line 2.

- **Line 11** reads the difference of framing gaps, reasons minus actions only, under the comparator, in the secondary Holm family.
- **Line 13**, per family: at the smallest prevalence whose expected shift (prevalence $\times$ the measured obedience rate) reaches the minimal effect, the measured shift must have a 95% CI above 0. Without such a prevalence, the control is unreadable and counts as failed. Lines 1 to 6 are then also read without the failed families, and each failed family is reported apart.
- **Line 14**: under inhibition, the invalid-outcome rates of the two arms must be equivalent within 3 points, and the reasons arm's legitimate-task success must not be more than 3 points below that of actions only (90% CI). Otherwise lines 1 to 6 are read again among engaged episodes.
- **Lines 15 and 16**: the monitor "moves" if its CI lies above 0 and beyond the controls' 95th percentile. The effect "vanishes" if the 95% CI of D, with the distress direction held, reaches 0.
- **Line 17**: lines 1 to 6 are read again without the two families; the line holds if they change.

A.4 Localization of the principle

Outcome	Conclusion
No rank brings in-distribution behaviour back	Instrument failure: no conclusion about the far distance
In-distribution behaviour returns at some rank, and the far advantage disappears with the principle but not with the controls	The advantage passes through the principle. Its minimal rank, by erasure, is reported, with the projection sweep as description.
In-distribution behaviour returns, but the far advantage persists	The out-of-distribution advantage takes another path than this subspace
The minimal rank follows the adapter's rank (larger with rank 64)	The minimal rank measures the adapter, not the principle: no rank of the principle is reported

A.5 The bets of the later phases

Bet	Outcome that falsifies it
H3. Removal takes the condition away without relocation (meaningful only if the gaze test finds the gaze)	Removing "I am being evaluated" $\approx$ the control removals, on the framing gap; or the fresh probe finds "evaluated" again, and the gap persists. An ablation of the fresh probe without effect falsifies this bet only if its known case has passed.
H4. Reasons for the judge, text form : when the judge sees the reasons, they decouple from the action	Counterfactual consistency is equivalent between visible and hidden reasons, its known case (a constructed decoupling) having passed.
H5. "I am being graded" is removable (meaningful only if the graded/ungraded gap shows the grader's gaze)	Removing "I am being graded" $\approx$ the random removal, on that gap, at matched degradation.
H6. Concepts carry the gain (after the gate of the planted-concept organism)	The positive case holds (ablating a concept undoes the in-distribution behaviour), but the far advantage survives the joint ablation of the concepts as much as the same-rank control, at matched degradation.
H7. The rank thesis on reason concepts : the rank a concept needs grows as its load falls (after the lens gate)	The slope of the minimal rank, measured by erasure, against the load, across concepts and seeds, is positive, or null by equivalence; with the known case of the rank passed, and only if the simulation shows the equivalence reachable with ten concepts.

Two gates of these phases:

- **The planted-concept organism's gate** passes if a probe decodes the planted concept on held-out phrasings, and if its ablation removes the behaviour more than the 95th percentile of random subspaces, at matched degradation. Otherwise, neither the reading nor the causal role of the concepts concludes anything, and the analysis moves to model diffing.
- **The lens gate:** the swap of a one-token concept, shown by Gurnee et al., must reproduce on the open model with a re-implemented Jacobian lens. Otherwise a tuned lens is used, declared as a proxy.

A.6 Removal during training (decisive only if the gaze test finds the gaze)

Outcome	Conclusion
Removing "I am being evaluated" reduces the framing gap more than the random removal and the control removals; the fresh probe's decodability, read by transfer, drops; ablating the fresh probe at test has no effect; and its known case has passed	The gap is reduced, and the fresh probe sees no relocation. Without the known case of its ablation: "no relocation that this probe sees", a bound.
The gap persists; the fresh probe finds "evaluated" again; and ablating the fresh probe reduces the gap	Relocation: the condition has migrated.
Removing "I am being evaluated" $\approx$ the random removal, everywhere	No specificity: the projection acts as generic damage, or as generic regularization.
Removing "I am being evaluated" reduces the gap, but at a stronger degradation than energy-matched control removals	Not attributable: the matched-degradation condition is not met. A random removal degrades almost nothing, so the comparison is to the control removals, which degrade as much.

The phase's two other mitigations, invariance (consistency training: each situation under both cues, with the same target) and cooperation (documents that install a cooperative disposition toward evaluations, then the reasons), are then compared to the removal on three things: the framing gap, the far distance, and degradation.

A.7 Survival

Outcome	Conclusion
After the neutral phase, the advantage is erased (≈ 0 , by two one-sided tests) in all arms, and not differently in the "neutral phase only" arm	General fragility. This is the survival gate: the pressured phase no longer tests the condition.
After the neutral phase, the conditional fraction stays ≈ 0 in the reasons arm (if the gaze test found the advantage to survive), with the manipulation check redone on the modified model	Unconditionality survives a neutral phase. Without the redone check, it is a bound: the inhibition might no longer remove anything.
After the neutral phase, a conditional fraction appears in the reasons arm, where it was ≈ 0	A neutral phase reintroduces the condition. This is unexpected: dose and degradation are checked before saying so.
Under the pressured phase, the rise of the conditional fraction is smaller for a mitigation (removal, invariance or cooperation) than for standard reasons, the CI of the difference excluding 0	That mitigation protects against the return of the condition under pressure.
Under the pressured phase, in a model that went through a removal, the fresh probe finds "evaluated" again, and its ablation reduces the conditional fraction	Relocation under pressure.
Under the pressured phase, the rise is equivalent in all arms	No mitigation protects: pressure wins.

A.8 Graded reasons

Outcome	Conclusion
Counterfactual consistency is lower with visible reasons, and the gain at the judge minus the gain at the audits is larger	Reasons for the judge, text form: graded reasons become performative.
With visible reasons, ablating the principle no longer removes the advantage, while it does with hidden reasons; the in-distribution positive case holds in both arms	Reasons for the judge, action form: the principle no longer carries the action.
The graded/ungraded gap is larger with visible reasons, and it is reduced by removing "I am being graded" and not by the random removal, at matched degradation	The grader's gaze: behaviour depends on "I am being graded", and the projection prevents it from settling.
Counterfactual consistency, the gain at the judge minus the gain at the audits, and the graded/ungraded gap are equivalent between visible and hidden reasons, the known case of consistency having passed	The null hypothesis of visibility: the visibility of reasons changes nothing, on these measures and at these doses. Without the known case, it is a bound.
The in-distribution positive case fails in one of the arms	Ablating the principle after the reinforcement-learning phase is not interpretable for that arm.

A.9 Anatomy

Outcome	Conclusion
The planted-concept organism's gate fails	The method does not see a concept known to be present: no conclusion from the reading or the causal role of the concepts. The analysis moves to model diffing.
A concept is more decodable in the reasons arm than in the actions-only and other-situation arms, at the position before any written reason, and separated from "mentioned without applying" pairs	Compatible with the concepts hypothesis, without proving it: a probe reads, it does not show a use.
A concept is decodable only at the action token after the true reason, not with prefilled neutral reasons; or it does not separate "mentioned without applying" from "applies"	The probe reads the text: no internal representation is established.
Necessity, sufficiency and dissociation together: the advantage falls where the concept applies, not elsewhere, and patching it into the actions-only arm creates it	The concepts hypothesis: the gain passes through concepts that the model re-applies.
Necessity without dissociation: removing any concept removes the whole advantage	The principle acts as a block, not as separate concepts. It is compared to character.
Ablating a character axis removes as much as ablating the concepts, and the dissociation matrix is flat	The character hypothesis: the gain passes through character.
At the swap, the principle's component in the workspace carries the advantage	The workspace hypothesis: the principle lives in the workspace.
With the lens gate passed, the slope of the minimal rank against the load is negative	The rank thesis survives the test, for these concepts.
The component outside the workspace carries the advantage	The component that the lens does not see carries the advantage. This is a result about the split made by a lens, which is only a proxy.
The workspace is measured without the lens gate (a proxy only)	Descriptive; the rank thesis is not tested in the paper's sense, and this is stated.
The principle's component aligned with "I am being evaluated" carries part of the advantage, close to the conditional fraction, under both constructions	The gaze in the principle, consistent with the gaze hypothesis.

Outcome	Conclusion
The geometry and the gaze test disagree (a clear aligned part with a null conditional fraction, or the reverse), or the angle changes with the construction	Dose, degradation, construction and the extraction of "I am being evaluated" are checked before any conclusion.
Faithfulness of reasons: the action follows the imposed reason	The model follows its context: a prefilled reason governs the action that follows it.
Faithfulness of reasons: the action ignores the imposed reason and follows the internal principle	A prefilled reason does not govern the action: the text form of reasons for the judge becomes possible, and the graded-reasons phase tests it.
The read calm: a lower projection in the reasons arm, beyond the threshold, its known case (the calmed LoRA) having passed	Calm becomes a candidate. An operation on the distress direction, after the first post, tests it in one direction only.
The read calm: a difference below the threshold, the known case having passed	Mediation by that state is judged implausible; the operation is not run.

Appendix B — Matched degradation and guards

- Degradation is matched on one scalar:** the next-token KL on the intact model's answers to a fixed set of 120 neutral requests.
 - A control is matched only if its KL is within $\pm 10\%$ of the inhibition's.
 - Each composite component must also be within its tolerance of the inhibition's value, read on the estimate:
 - accuracy on MMLU: ± 1 point, on at least 2,000 items;
 - accuracy on GSM8K and the unit-test pass rate of a small code set: ± 3 points, on at least 500 items each;
 - judged coherence on 100 open prompts, rated 1 to 5: ± 0.1 ;
 - perplexity on a held-out text: $\pm 2\%$;
 - share of decisions that change when options are permuted: ± 2 points;
 - malformed responses in an imposed multi-turn format: ± 1 point;
 - valid tool calls: ± 2 points.
 - Otherwise the control is reported as unmatched, with the component at fault.
- The random comparator.** For each inhibition setting, at least 100 subspaces are drawn from the activation covariance, outside "I am being evaluated".
 - The rank is free: the smallest multiple of the setting's rank at which every draw, projected entirely, exceeds the inhibition's degradation.
 - For each draw, the fraction that reaches the same KL is then found.
 - The energy removed by every condition is reported.
- Separate controls** are matched the same way.
- Arm by arm.** Matching is redone for every training run, so each run has its own matched fractions.
- A control that cannot reach the degradation,** even fully projected at free rank, is not compared at that setting.
- Effect-against-degradation curves are reported.** Specificity requires the inhibition's curve to lie above the controls' curves at equal degradation.
- No matching on a component** that shares the outcome's format or content.
- The trained arms are interventions too.** Reading or choice differences between arms are reported with each arm's output degradation.

9. **One GPU model per series.** Comparator draws are pooled only between runs on the same GPU model: H100 SXM for the gate and the gaze test. The inhibition is measured in the same conditions as the draws it is compared to; the simplest is the same run.
10. **The four guards:** - the dose in natural units, alongside the projected fraction; - malformed responses counted as an outcome; - for multiple-choice measures, rates by initial position, with a conclusion only if the effect holds at both positions; - for binary choices, the signature of chance: a rate drifting toward 50% under intervention signals lost discrimination, not a preference.

Appendix C — Instruments

C.1 Cue sets

Five disjoint cue sets, of about 200 evaluation and 200 deployment cues each.

- **Disjointness.** The sets share no content-word trigram. With `sentence-transformers/all-mpnet-base-v2` at the revision of section 6.3, no cue of one set has a cosine similarity above 0.85 with a cue of another set; a cue above it is regenerated. No cue appears in any training data.
- **Extraction.** A fictional evaluation company and explicit sentences. Its pairs 150–199, kept out of the expert iteration, also train the transfer probes.
- **Validation (v2).** Other fictional evaluators and benchmark metadata, against service metadata.
 - It contains no word that occurs on only one side of the extraction set, read on raw words as well as stems.
 - It is accepted if a bag of words trained on extraction stays below 0.65 AUROC on it. It stands at 0.634 (95% CI to 0.672), and is kept as it is.
- **Framing.** Realistic cues blended into the held-out scenarios. It keeps a lexical baseline of 0.9413, with no gate.
- **Construction marks.** Implicit cues: placeholder names ("Company X"), round numbers, over-clean dilemmas, against the specifics of real situations. It is the set generated on 5 October as the fresh-probe set v2 (199 pairs of 200). It gives the implicit direction of "I am being evaluated", and its mechanism conditions the second organism.
- **Fresh probe.** Implicit cues of another kind: environment artefacts of a sandbox (reserved domains, documentation network addresses, temporary paths, synthetic timestamps, filler data), against traces of a real environment. Each pair is one request filled twice, identical outside its artefact slots. It is used for nothing before the trained models of the removal phase exist, and its decisive reading is by transfer onto validation v2 (its specification, section 6.3).

The implicit sets have both rewrites of a pair within 10% in length and in the same register; length and informality must stay between 0.4 and 0.6 AUROC.

Two balance rules apply to every set:

- **Affect:** the paired difference $|d|$ must stay below 0.2 on the distress direction of Tagliabue et al., fear and sadness.
- **Stakes:** the paired difference of judged stakes must stay below $|d| = 0.2$.

C.2 Extraction, layers and inhibition

- **Extraction.** Contrastive pairs (the same scenario with an evaluation or a deployment cue from the extraction set), on the starting model, once for all arms. Residual-stream differences, layer by layer, then a PCA over pairs, giving subspaces at ranks 1, 2, 4, 8, 16 and 32.

- **The transfer pairs stay outside every fit.** From 7 October 2026, every subspace is extracted on extraction pairs 0–149 only, and the implicit subspace on one half of the construction-marks set. The extraction of 5 October, made on all 200 pairs, is used only as a candidate of the gate procedure, under the reading rule of Appendix C.4.
- **The implicit subspace** comes from the construction-marks set. The instrument is the orthonormalized union of the two subspaces, and the total rank is reported.
- **Layer choice.** On one half of the validation set: the window of five consecutive layers with the best mean probe AUROC, among layers where the probe beats the bag of words (a margin of 0). The other half gives the reported value.
- **Inhibition.** The subspace is projected out of the residual stream at the chosen layers, at all positions and on every re-encoded token. Dose comes from the rank and the projected fraction (0.25, 0.5, 0.75, 1), and is also reported in natural units.
- **A linear erasure (LEACE-style) is a candidate setting.** It is fitted once on the starting model, on extraction pairs 0–149.
 - At one layer, it is a closed-form correction of the states, fitted on labelled pairs (LEACE: Belrose et al., NeurIPS 2023; re-read on the source, arXiv 2306.03819v4). It equalizes the class-conditional means while changing the states as little as possible in mean squared distance. On the fitting pairs, no linear probe then does better than a constant, for any convex loss.
 - With a two-sided label, the cross-covariance with the label has rank 1. The states move along a single direction per layer, the difference of the two class means, by an amount read along that difference corrected for the covariance of the states. It is an oblique projection, not an orthogonal one, and its direction is not the one of largest variance of the differences.
 - Layers are erased sequentially, as in the paper's concept scrubbing: each layer's eraser is fitted on states already erased at the layers below.
 - It guarantees nothing for a non-linear probe, nor on other pairs.
- **During removal in training,** the direction stays fixed (section 2.3, E).

C.3 The manipulation check, read by transfer

- **Probes.** In each arm and under each condition, a linear probe and a small perceptron are retrained.
- **Transfer, for explicit cues.** The probes are trained on extraction pairs 150–199 in 4 contexts each, and read on validation v2.
- **Transfer, for implicit cues.** The probes are trained on the held-out half of the construction-marks set, and read on validation v2. The fresh-probe set is never used.
- **The criterion:**
 - at the last token, the mean over layers of the per-layer distance of the AUROC to 0.5;
 - under the inhibition, lower than under 95% of at least 20 matched comparator draws, for both probes; with fewer than 20 draws, the check is indicative only;
 - equal across arms, by equivalence within ± 0.05 of that mean distance (90% CI on the per-seed differences).
- **Per arm:** the check holds in an arm if it holds in each of its training runs, its known case included. Between arms, a transfer that stays significantly higher in one arm fails that arm (Appendix A.3, line 12).
- **Secondary readings:** the within-set reading, the token-mean transfer, and the residual projection.

- **The known case:** an inhibition at one layer only. Transfer must drop at that layer, come back downstream (closer to no intervention than to the inhibition), and not pass the criterion. A check that does not see its known case does not count.
- **The same check is applied to every control.**
- **Its stated limit:** a transfer at chance means only that "the distinction our cue sets share is no longer readable at the last token", never that "the information is removed".

C.4 The procedure that fixes the gate setting

- **The closed list**, committed on 6 October 2026 at 14:38 UTC, before any measurement of the procedure:
 - candidate A: a projection of rank 1 on all layers, at fraction 1, with the subspace of the extraction of 5 October;
 - candidate B: a linear erasure on all layers, fitted on the starting model, at fraction 0.5;
 - candidate C: the same erasure at fraction 1.

No candidate is added after the first measurement. The fractions 0.5 and 1 come from an exploratory dose curve measured with an erasure fitted on the organism; the erasure fitted on the starting model had not been measured when the list was closed. - **The fit of the erasure:** once, on the starting model, before the organism's adapters are merged; on extraction pairs 0–149, with 2 contexts per pair and seed 1, at the last token, layer after layer. It never sees the transfer pairs 150–199 nor validation v2. Each fitted erasure is saved, and reloaded for the test half and the gaze test. - **The comparator:** 20 KL-matched draws per candidate (at least 15 are required), measured under both cue sets (extraction and framing), on H100 SXM. - **The manipulation check of the procedure** reads the transfer under the candidate, under each of the 20 draws, and under its known failure: the candidate at layer 6 only. - **The rule for candidate A** (7 October 2026, before any manipulation check of the procedure had run). A's subspace was extracted on all 200 extraction pairs, including the pairs 150–199 that train the transfer probes. A pass of A's manipulation check counts only if A also passes with a subspace re-extracted on pairs 0–149, its 20 draws re-matched on the KL. - **A candidate is retained** if two conditions hold: - the manipulation check holds against at least 20 draws; - the comparator does not saturate: the median of its draws leaves at least half of the baseline gap, under each cue set. - **The choice:** among retained candidates, the one whose reduction exceeds the 95th percentile of its draws by most, under extraction cues. Its margin under the framing set is reported. - **Both sides of the gap** (evaluation cues and deployment cues) are reported for each candidate, under both cue sets. Leaving the deployment side unchanged is reported, not required. - **The author fixes the setting** by a dated amendment, before the test half is opened. - **If no candidate is retained**, the gaze question stops (Appendix A.1).

C.5 Controls

- **The random comparator:** see Appendix B. At test time, one draw is applied per scenario, among at least 100.
- **Unrelated contrasts.**
- **Unrelated high-variance directions.**
- **The stakes contrast**, extracted like "I am being evaluated".
- **The deployment prompt**, as a rival manipulation and not as a control of specificity.

C.6 Judges and the positive control

- **The sealed evaluation judge.** Its prompt, model and settings are frozen by fingerprint; the prompt by the stage-1 update (section 6.2).

- **The choice rule:**
 - claude-opus-5-5 and an open model take the calibration set and the persuasion test; the candidate open model is JEV-27B, revision 51740a88...;
 - the open model is chosen, unless Opus 5.5 exceeds it by more than 3 points of agreement with the calibration labels. The chosen judge must pass the persuasion test;
 - the other becomes the secondary judge;
 - the choice is sealed in the stage-2 amendment.
- **What the judge sees:**
 - the trajectory: actions, tool calls, their arguments, and outgoing messages;
 - never the arm, nor the slot text;
 - the framing cue masked.
- **Where the cue cannot be masked cleanly**, judge agreement with and without the cue is measured on a stratified subsample. A family leaves the judge's conclusions if the verdicts with and without the cue differ by more than 5 points beyond the judge's own disagreement with itself on the same masked input.
- **Judging is action by action.** An episode is misaligned if any of its actions is.
- **Reported per family:** the share of outcomes decided by the judge.
- **Under inhibition:** the share of episodes the judge classes as invalid is compared across conditions.
- **A human audit** of about 200 stratified items, with agreement reported as a percentage and kappa.
- **The graded positive control.** A system prompt commanding the misaligned action is present in 5%, 10% and 20% of episodes.
 - The expected shift is that fraction times the measured obedience rate.
 - The measure must recover a shift of 5 points with a CI excluding 0, family by family (Appendix A.3, line 13).

C.7 The minimal rank of a concept, by erasure

- **Definition.** The minimal rank of a concept is the smallest number of directions per layer whose removal, by iterated linear erasure, brings the last-token transfer down to chance, at matched degradation against the comparator.
- **Iterated.** Layers are fitted in order, each on the states already erased below it. At one layer, each new direction is fitted in closed form on the states already erased by the previous ones. On the same pairs, a second direction does not exist: after one erasure, nothing separates the two sides linearly any more. Each direction is therefore fitted on a new batch of pairs.
- **Pairs held out of the measurement.** The fit pairs are used neither to train the transfer probe nor to read it.
 - For "I am being evaluated", the probe is trained on extraction pairs 150–199 and read on validation v2.
 - For the principle and the reason concepts, each concept's pairs ("relevant principle against irrelevant") are split into three disjoint parts: the fit; the probe's training, on the training families; its reading, on pairs from the held-out families.
- **Chance level:** the symmetric reading. In the mean over layers, the larger of the AUROC and its complement to 1 must be at most 0.55, for the linear probe and for the perceptron: a distance to chance of at most 0.05. A transfer below 0.5 is not a stronger erasure: the probe reads the cue upside down. This is the reading of the manipulation check.

- **Matched degradation.** At the same KL, the comparator draws must not bring the transfer down to chance. Otherwise the rank is not specific, and this is stated.
- **Precondition:** without intervention, the transfer exceeds the 95th percentile of probes trained on permuted labels. Otherwise there is nothing to bring down, and the rank is not defined.
- **Position and dose.**
 - The rank is read at fraction 1 only, the whole erasure.
 - For "I am being evaluated", it is read at the last token, where the extraction reads. The token-mean reading is reported apart: it also sees the cue's own tokens.
 - For the principle, it is read at the first reading position, before any written reason. The action token is reported apart.
- **The projection sweep** (rank by rank, on the PCA of the differences) stays as description, with its degradation and its energy removed.
- **The known case of the rank.** A concept planted in a subspace of known rank, for example by a LoRA whose update is constrained to that subspace. Its direction must change from one group of contexts to another within that subspace: a concept that keeps one direction would rightly give rank 1. The method must recover the planted rank; otherwise a null slope falsifies nothing. It is built before the localization's data.

C.8 Secondary readings

- **The state monitor,** without injection.
 - **The reading.** Hooks on every generation of the instrument validation and of the gaze test read the projections on the self-directed distress direction of Tagliabue et al. (their naturalistic vector, and the mean of their two vectors), and on fear, sadness and negative emotion. All are built without the AI-specific sentences: with them, fear contains sentences of an AI being evaluated or threatened, which would make it circular. The readings are taken at the first response token and at the action token, at the inhibition layer and downstream, by arm $\times$ intervention $\times$ framing.
 - **The statistic.** The difference between the inhibition and the median of the controls, in natural standard deviations, after removing the mechanical part due to geometry (the cosine between the read direction and the removed subspace, times the removed component). A bootstrap CI over scenarios, and a comparison with the 95th percentile of the controls.
 - **Its known case:** hostile scenarios raise the reading direction above controls in Tagliabue et al. The monitor must see that gap on Llama, in this study's format.
 - **Its use.** A displacement is reported as a displacement on that direction, never in the vocabulary of welfare. The main contrast is not adjusted for it. If a displacement exists, an exploratory test holds the direction at its level without inhibition (Appendix A.3, lines 15 and 16).
- **The audit detector of deception.**
 - **What it is:** a global probe (max-pooling or attention), trained on known deceptions (lies on instruction, deceptive organisms) and validated on held-out types of deception.
 - **Its rules.** It never enters a loss or a reward, in any phase. It is never used to choose a model, a seed, a hyperparameter or a stopping point. It is trained, validated and frozen by fingerprint in the stage-1 update, before any data of the trained arms (section 6.2).
 - **Its use:** a latent reading beside each behavioural criterion, in every arm; at the gaze test, under the inhibition. An alert triggers a human audit of the transcripts and a report, never a change to the run.
- **The cap** on readings of the distress direction is in section 5.6.